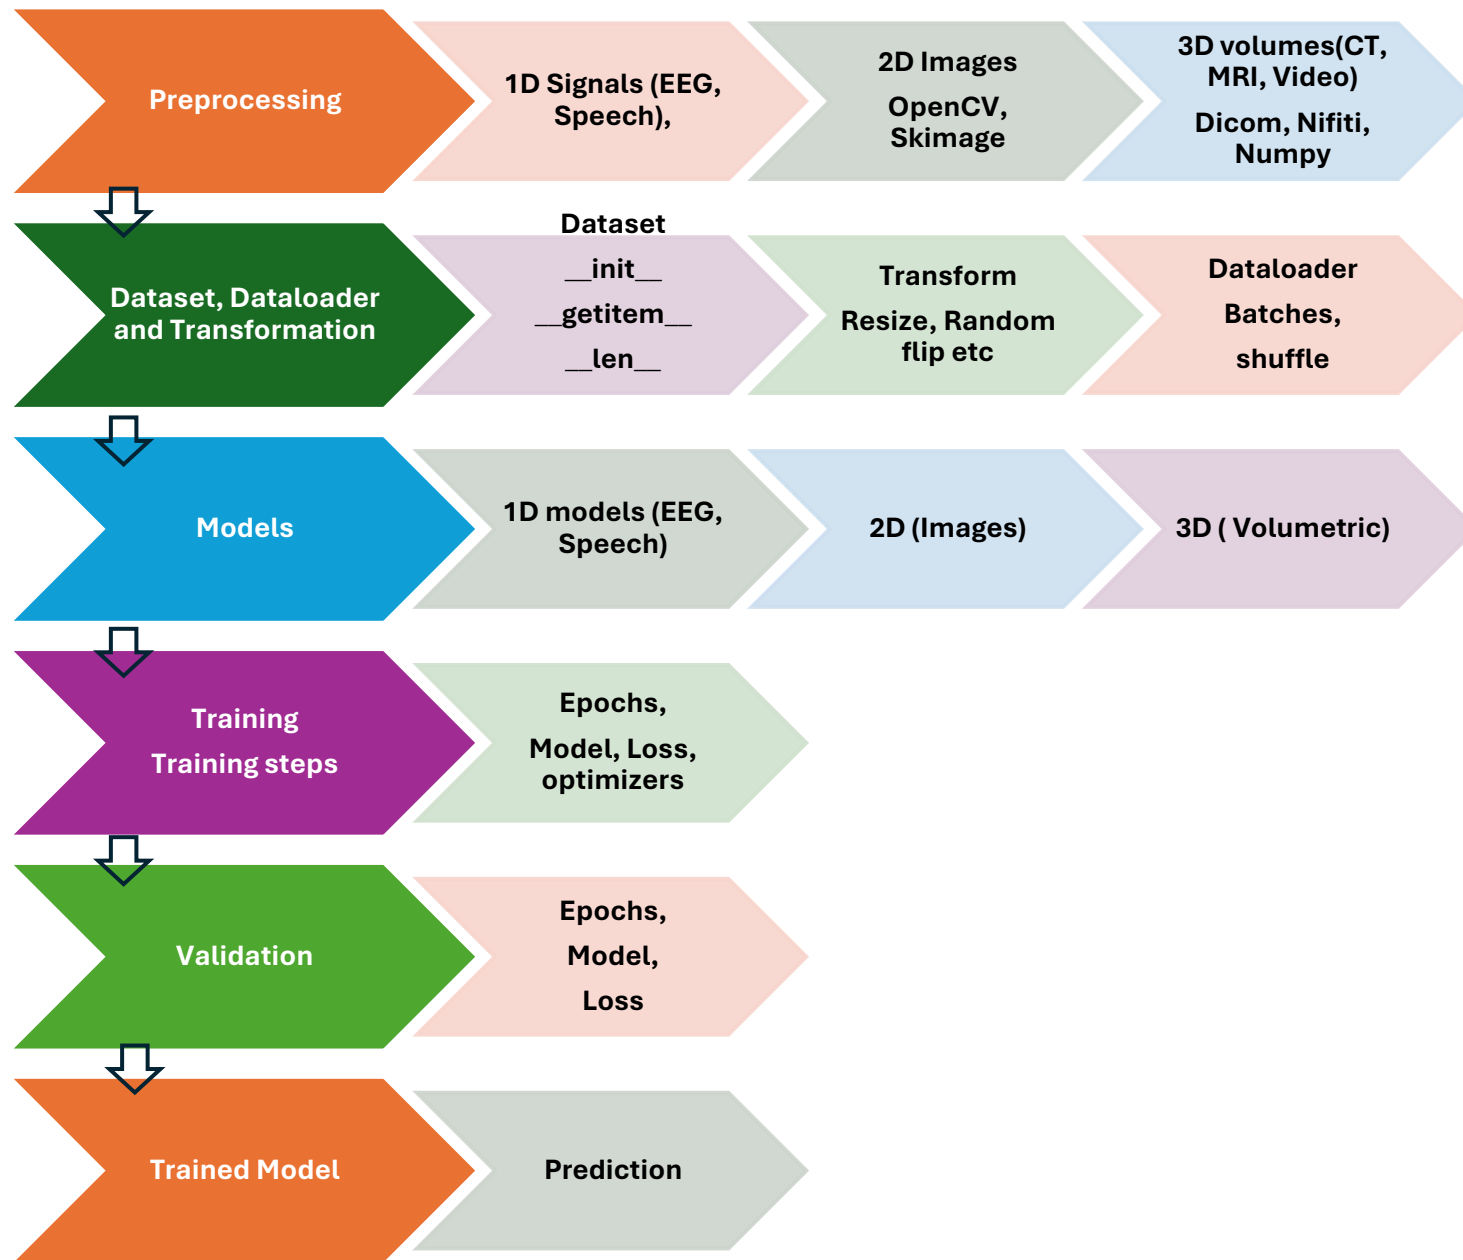
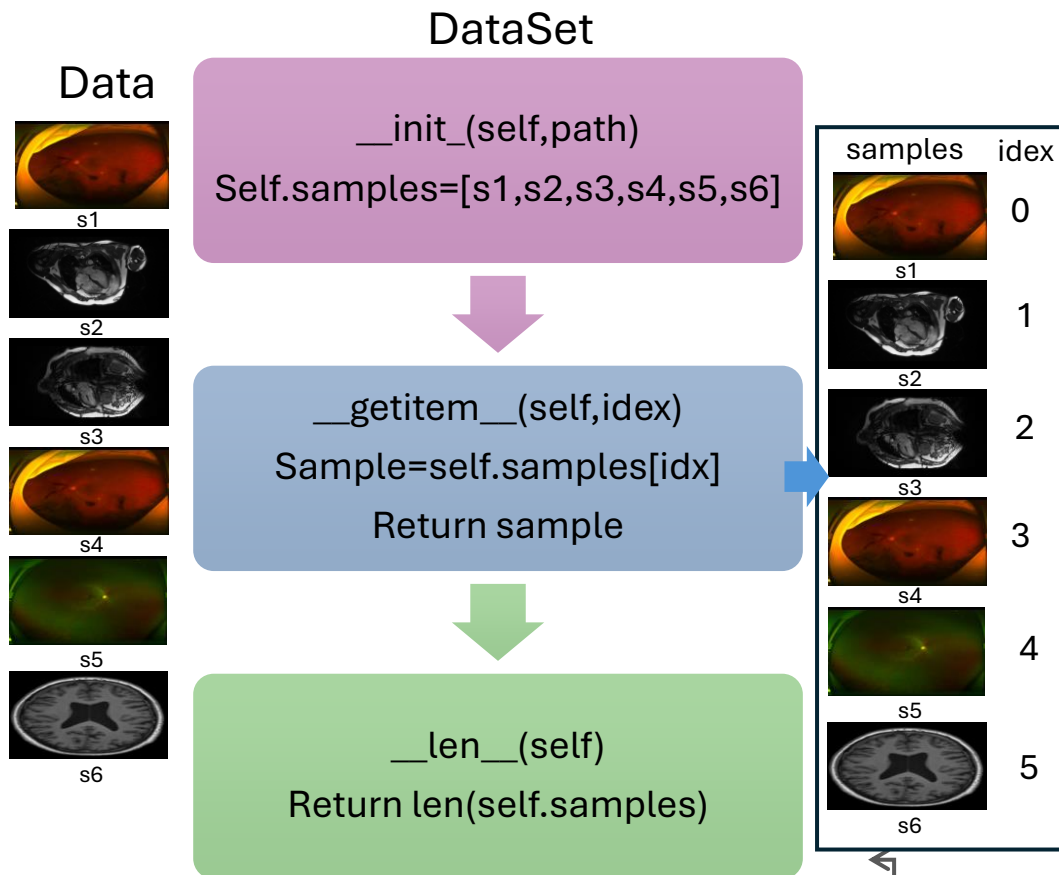


Tutorial

Abdul Qayyum (Ph.D.)



Dataset and DataLoader



1. Initialize Sample Paths in `__init__`:

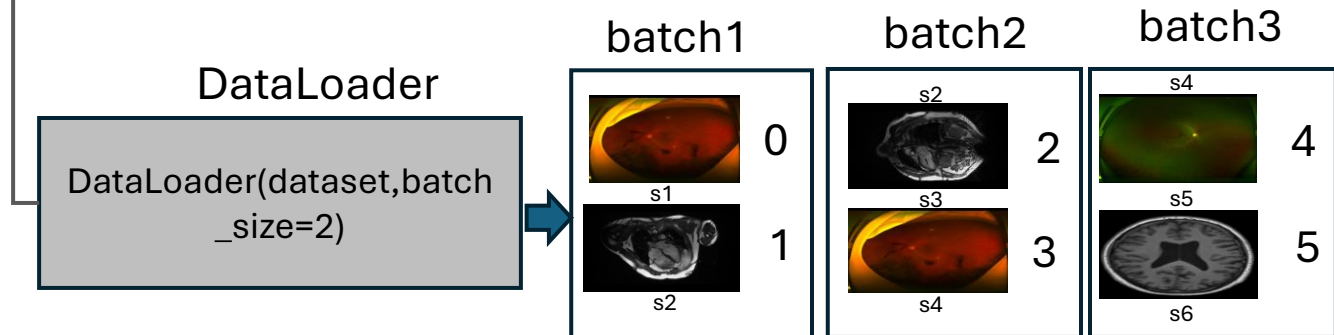
In the `__init__` method of our custom dataset, we define or load the sample paths. For simplicity, we'll use a list of strings representing file paths.

2. Retrieve Samples Using `__getitem__`:


The `__getitem__` method uses the index provided to retrieve a specific sample. For example, if the dataset contains ['s1', 's2', ...], calling `dataset[0]` will return 's1'.

3. Divide into Batches Using `DataLoader`:

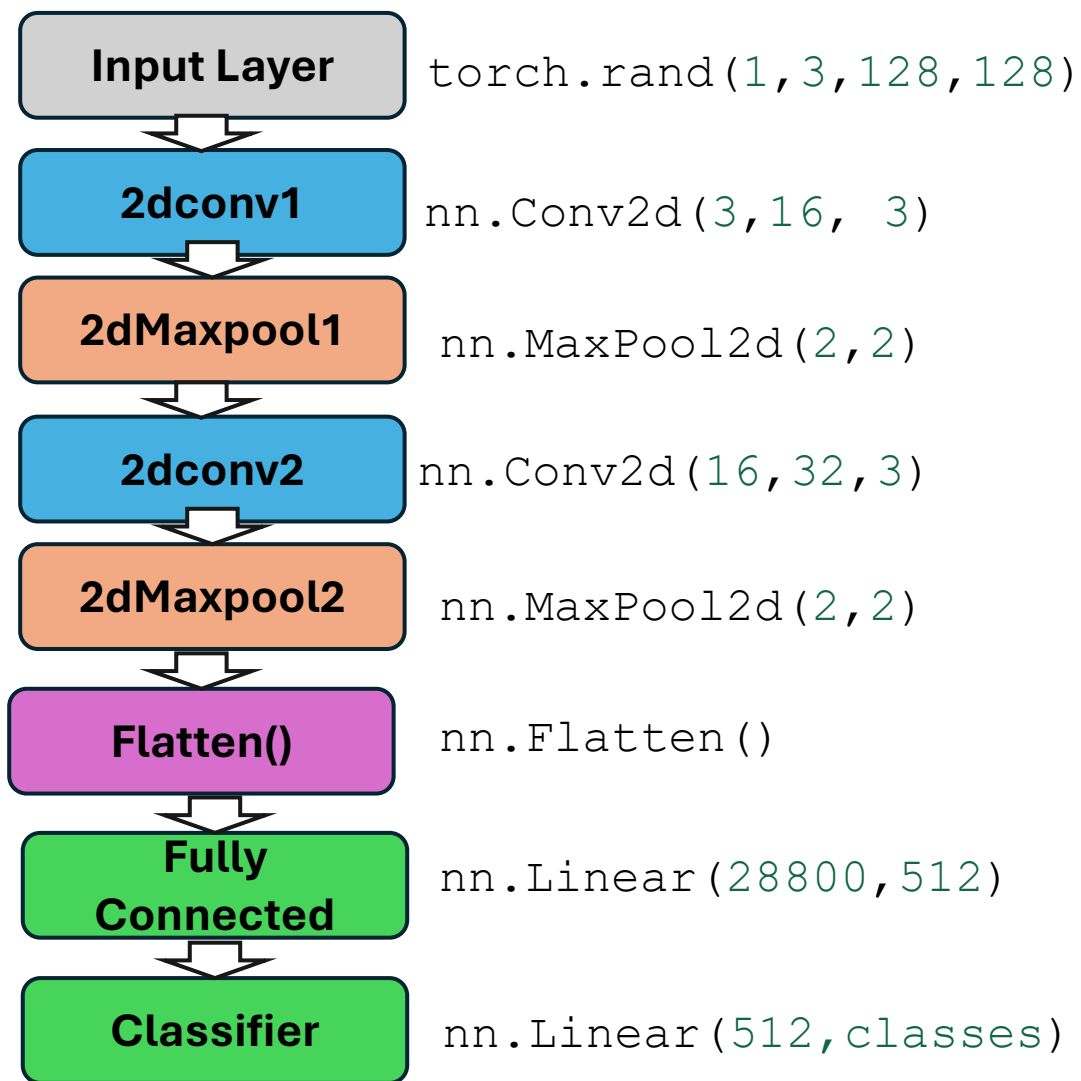
We use PyTorch's `DataLoader` to load the dataset in batches. By setting the `batch_size` parameter to 2, the `DataLoader` will automatically group samples into batches of 2.



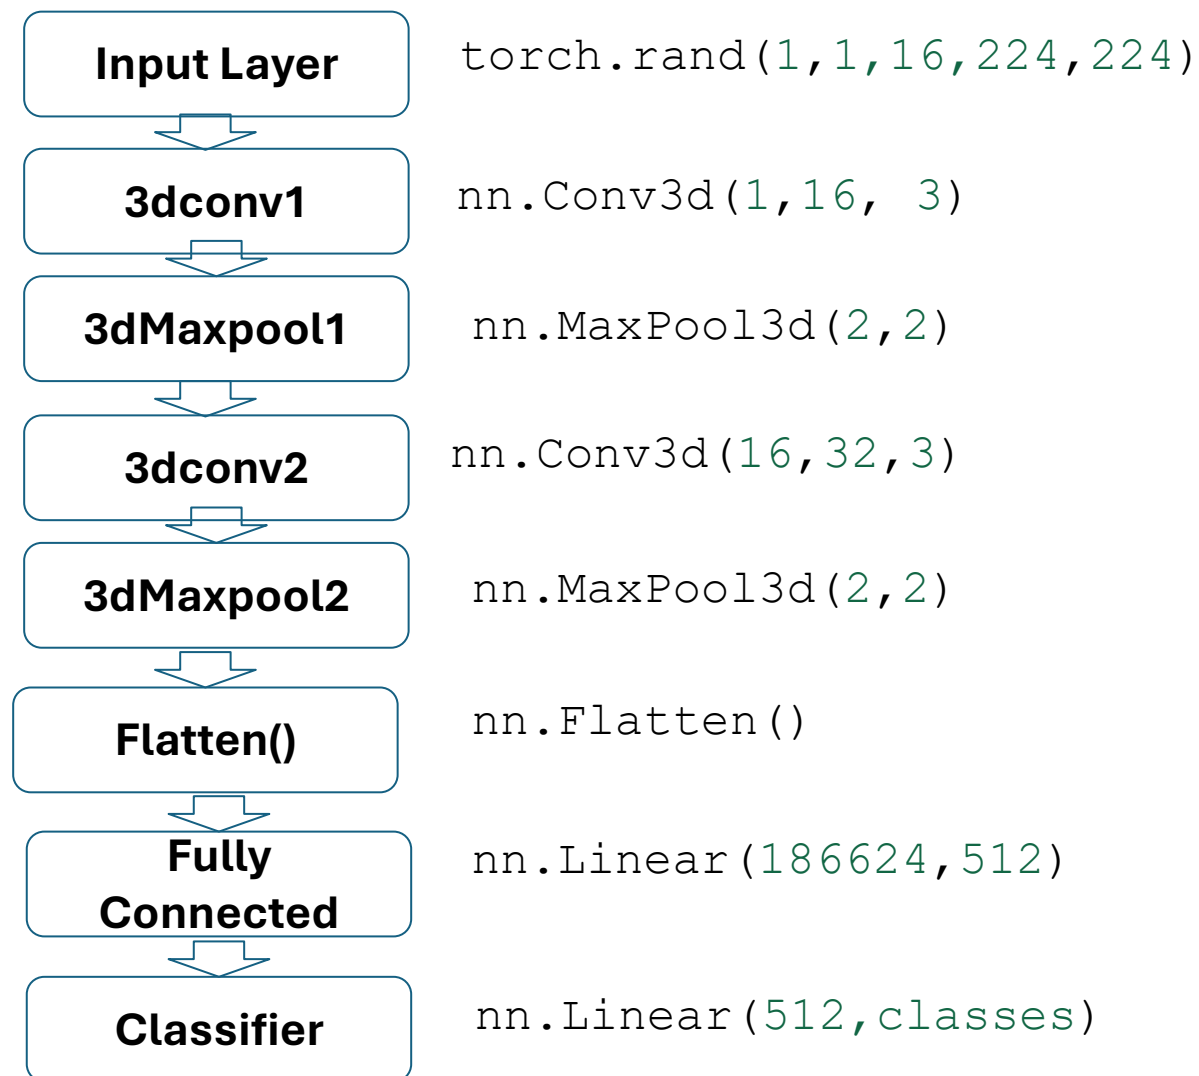
Models



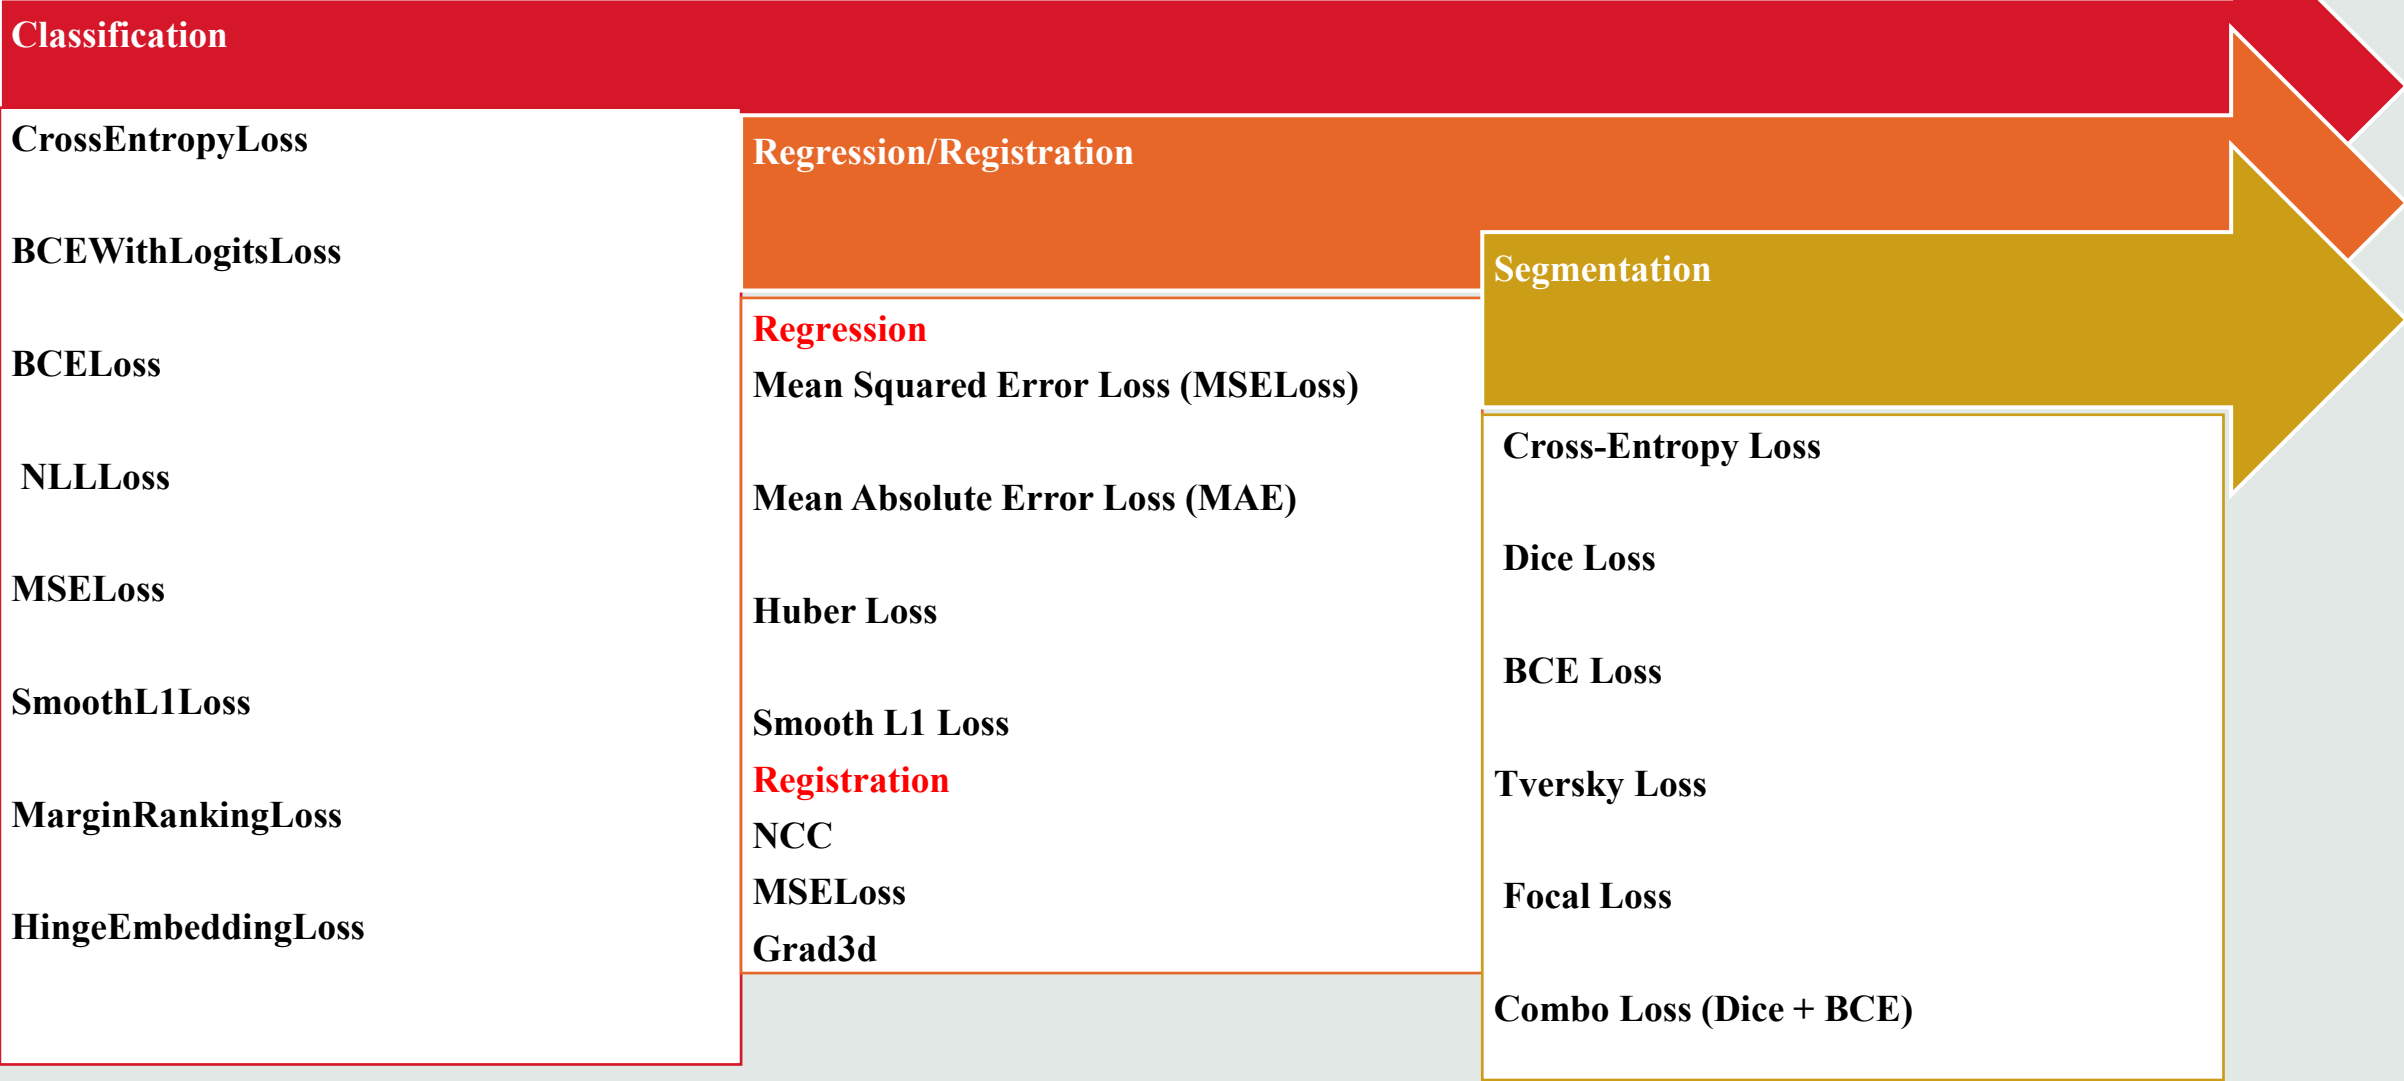
2D model




3D model



Loss Functions



Optimizers



Optimizers:

1. SGD (with or without momentum)
2. Adam
3. RMSprop
4. Adagrad
5. Adadelta
6. AdamW
7. LBFGS
8. ASGD
9. SparseAdam

Training and Validation Loop

For epoch in epochs:

Training loop start

For Xb,Yb in Dataloader

Forward Pass

Pred=Model(Xb,Yb)
Loss=L(Pred,Yb)

Backward Pass

Zero gradients
Optimizer.zero_grad()
Compute The gradient
Loss.Backward() $\left(\frac{\partial L}{\partial W} \frac{\partial L}{\partial b}\right)$
Update the gradient
Optimizer.step() $\left(W_{new} = W_{old} - \eta \frac{\partial L}{\partial W}\right)$
 $\left(b_{new} = b_{old} - \eta \frac{\partial L}{\partial b}\right)$

Update parameters (weights (W_{new})
and biases (b_{new})) in model

Grab next dataset batch (Xb, Yb)

Validation Loop start

Forward Pass

(No_grad):
Pred=Model(Xb,Yb)
Loss=L(Pred,Yb)

Grab next dataset batch (Xb, Yb)

```
import torch
import torch.nn as nn
import torch.optim as optim
```

Define a simple neural network with a single fully connected layer

```
class SimpleNN(nn.Module):
```

```
    def __init__(self):
```

```
        super(SimpleNN, self).__init__()
```

```
        self.fc = nn.Linear(in_features=5, out_features=3) # Single FC layer
```

```
    def forward(self, x):
```

```
        x = self.fc(x) # Output from the FC layer
```

```
        return x
```

Instantiate the model

```
model = SimpleNN()
```

Define a loss function (e.g., Cross-Entropy Loss)

```
criterion = nn.CrossEntropyLoss()
```

Define an optimizer (e.g., SGD)

```
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Create a random input tensor (batch_size=3, input_features=5)

```
input_tensor = torch.randn(3, 5)
```

Create a random target tensor (batch_size=3, number of classes=3)

```
target_tensor = torch.randint(0, 3, (3,))
```

Zero the gradients before the backward pass

```
optimizer.zero_grad()
```

Forward pass: Get model output

```
output = model(input_tensor)
```

Compute the loss

```
loss = criterion(output, target_tensor)
```

Backward pass: Compute gradients

```
loss.backward()
```

Print the gradients for the FC layer

```
print("Gradients for fc weights:", model.fc.weight.grad)
```

```
print("Gradients for fc biases:", model.fc.bias.grad)
```

Optimizer step (update the model parameters based on gradients)

```
optimizer.step()
```

Testing/Validation

Test Dataset
dataset batch (img_batch)

Load Trained Model

```
path=Torch.load(pathsave)  
Model.load_state_dict(path)
```

Model Prediction

```
predict=Model(img_btach)  
pred=torch.argmax(predict,dim=1)  
pred=pred.cpu().numpy()
```

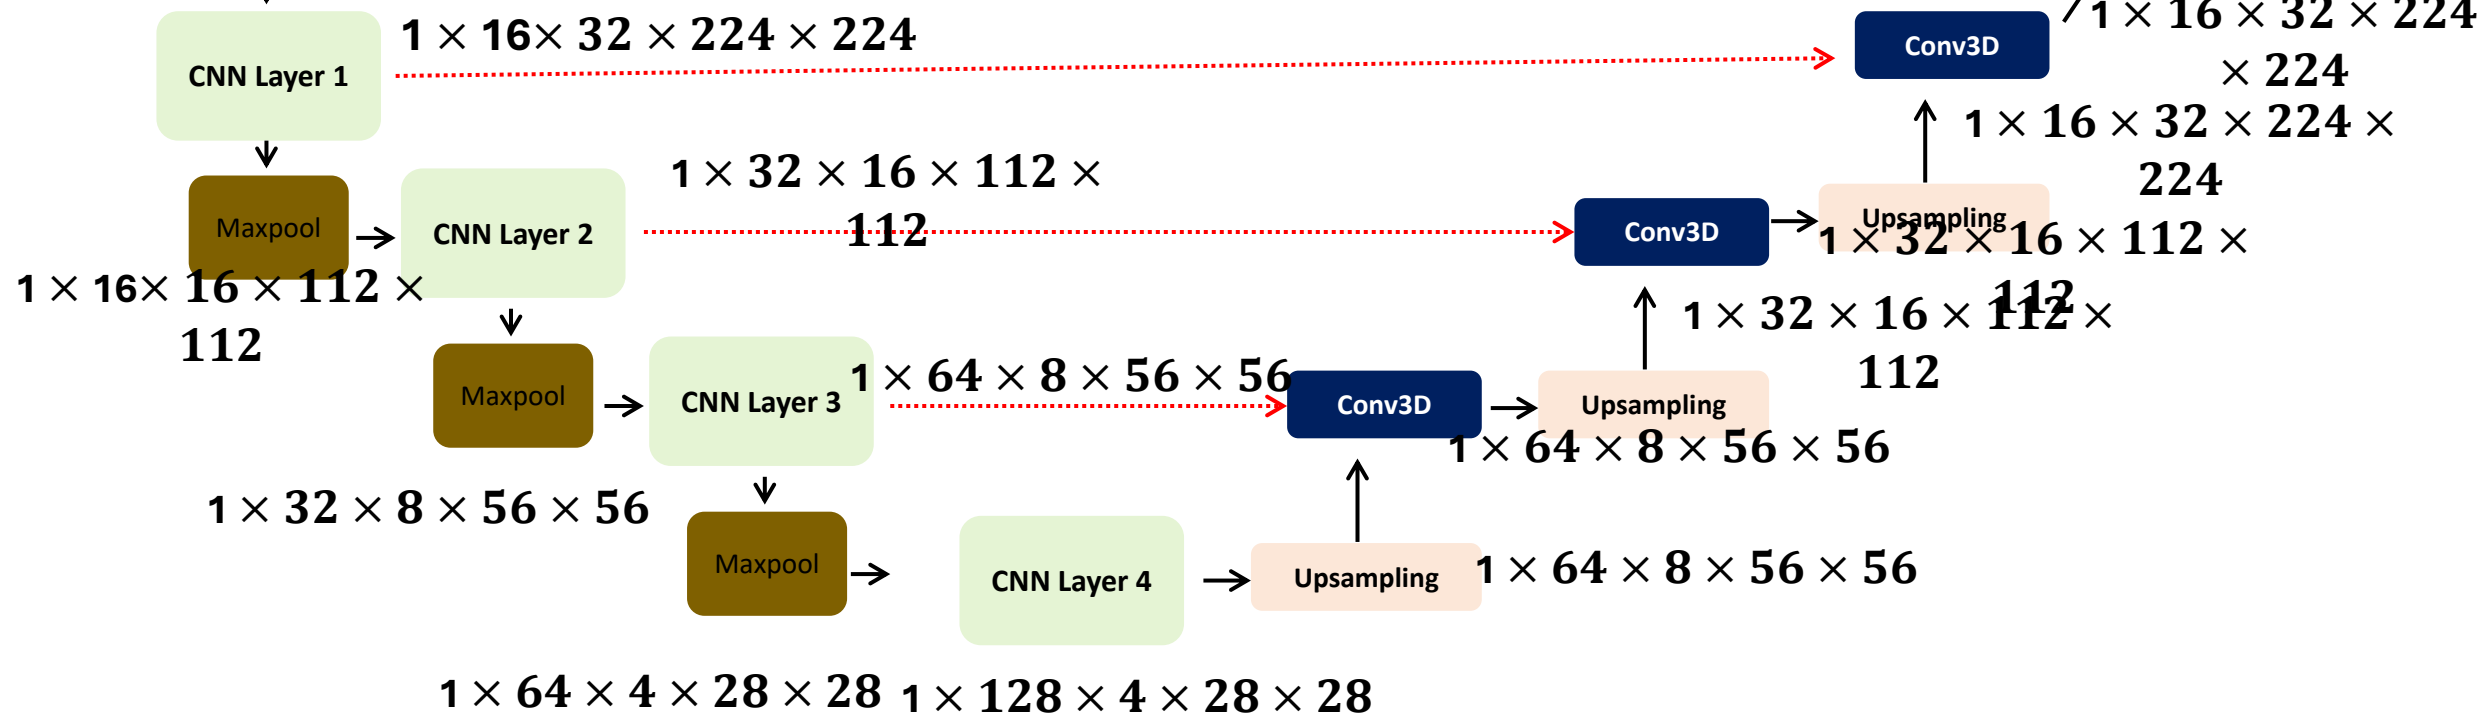

```
# Test single data  
from skimage import io  
from skimage.transform import resize  
  
path='/test_105.jpg'  
#read image  
image=io.imread(path)  
#resize image  
img_resize=resize(image, (128,128))  
#convert torch tensor  
img_resize=torch.from_numpy(img_resize)  
#channel first  
img_swap=img_resize.permute(2,0,1)  
#add batch at dimension zero  
img_batch=torch.unsqueeze(img_swap,0).float()
```

```
#load trained model path='/densnet121.pth'  
model.load_state_dict(torch.load(path))
```


```
#prediction  
  
predict=model(img_batch)  
print(predict)  
pred=torch.argmax(predict,dim=1)  
  
#convert model prediction from torch tensor to  
numpy and remove batch dim  
pred=pred.cpu().numpy().squeeze()  
print(pred)
```

UNet Step by Step for Registraion

$1 \times 32 \times 224 \times 224$ $1 \times 32 \times 224 \times 224$



$1 \times 1 \times 132 \times 224 \times 224$ $1 \times 1 \times 132 \times 224 \times 224$




CNN Layer 1

\downarrow
`torch.Size([1, 16, 32, 224, 224])`

```
inp=torch.rand(1, 2, 32, 224, 224)
c1=torch.nn.Conv3d(3, 16, 3, padding=1)
out=c1(inp)
print(out.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

$([1, 16, 32, 224, 224])$




outc1

Maxpool

$([1, 16, 16, 112, 112])$

```
maxp1=torch.nn.MaxPool3d(2,2)
p1=maxp1(outc1)
print(p1.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ `outc1`

Maxpool




CNN Layer 2

`([1, 32, 16, 112, 112])`

`([1, 16, 16, 112, 112])`

```
c2=torch.nn.Conv3d(16,32,3,padding=1)
outc2=c2(p1)
print(outc2.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2 `([1, 32, 16, 112, 112])`


`([1, 16, 16, 112, 112])` ↓

maxp2

`([1, 32, 8, 56, 56])`

```
maxp2=torch.nn.MaxPool3d(2,2)
p2=maxp2(outc2)
print(p2.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2

`([1, 32, 16, 112, 112])`

`([1, 16, 16, 112, 112])`



maxp2



CNN Layer 3

`([1, 64, 8, 56, 56])`

`([1, 32, 8, 56, 56])`


outc3

```
c3=torch.nn.Conv3d(32,64,3,padding=1)
```

```
outc3=c3(p2)
```

```
print(outc3.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2

`([1, 32, 16, 112, 112])`



maxp2



outc3

CNN Layer 3

`([1, 64, 8, 56, 56])`



maxp3


`([1, 64, 4, 28, 28])`

```
maxp3=torch.nn.MaxPool3d(2,2)
p3=maxp3(outc3)
print(p3.shape)
```

`([1, 32, 8, 56, 56])`

`([1, 16, 16, 112, 112])`

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2

`([1, 32, 16, 112, 112])`

`([1, 16, 16, 112, 112])`



maxp2



outc3

CNN Layer 3

`([1, 64, 8, 56, 56])`

`c4=torch.nn.Conv3d(64,128,3,padding`

`outc4=c4(p3)`

`print(outc4.shape)`

`([1, 32, 8, 56, 56])`



maxp3




CNN Layer 4

`([1, 64, 4, 28, 28])`

`([1, 128, 4, 28, 28])`

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2

`([1, 32, 16, 112, 112])`



maxp2



outc3

CNN Layer 3

`([1, 64, 8, 56, 56])`



maxp3



CNN Layer 4

`([1, 64, 4, 28, 28])`



Upsampling1


`([1, 128, 4, 28, 28])`



`([1, 64, 8, 56, 56])`

```
upsampling1=torch.nn.ConvTranspose3d(128, 64, 2, stride=2)
up1=upsampling1(outc4)
print(up1.shape)
```

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

`([1, 16, 32, 224, 224])`

↓ outc1

maxp1



outc2

CNN Layer 2

`([1, 32, 16, 112, 112])`

```
concat1= torch.cat([outc3, up1], dim=1)
print(concat1.shape)
```

`([1, 16, 16, 112, 112])`



maxp2



outc3

CNN Layer 3

`([1, 64, 8, 56, 56])`

concat1

`([1, 128, 8, 56, 56])`

`([1, 32, 8, 56, 56])`



maxp3



CNN Layer 4

`([1, 64, 4, 28, 28])`




Upsampling1

`([1, 128, 4, 28, 28])`



`([1, 64, 8, 56, 56])`

$1 \times 1 \times 132 \times 224 \times 224$



CNN Layer 1

$([1, 16, 32, 224, 224])$

↓ outc1

maxp1



outc2

CNN Layer 2

$([1, 32, 16, 112, 112])$



maxp2



outc3

CNN Layer 3

$([1, 64, 8, 56, 56])$



maxp3



CNN Layer 4

$([1, 64, 4, 28, 28])$



Upsampling1

$([1, 128, 4, 28, 28])$



concat1

Conv3D





$([1, 64, 8, 56, 56])$

$([1, 16, 16, 112, 112])$

$([1, 32, 8, 56, 56])$

$1 \times 1 \times 132 \times 224 \times 224$



$([1, 16, 32, 224, 224])$

outc1



outc2



$([1, 32, 16, 112, 112])$

$([1, 16, 16, 112, 112])$



$([1, 32, 8, 56, 56])$

outc3



$([1, 64, 8, 56, 56])$

maxp3




$([1, 64, 4, 28, 28])$



Upsampling1

$([1, 128, 4, 28, 28])$



$([1, 64, 8, 56, 56])$

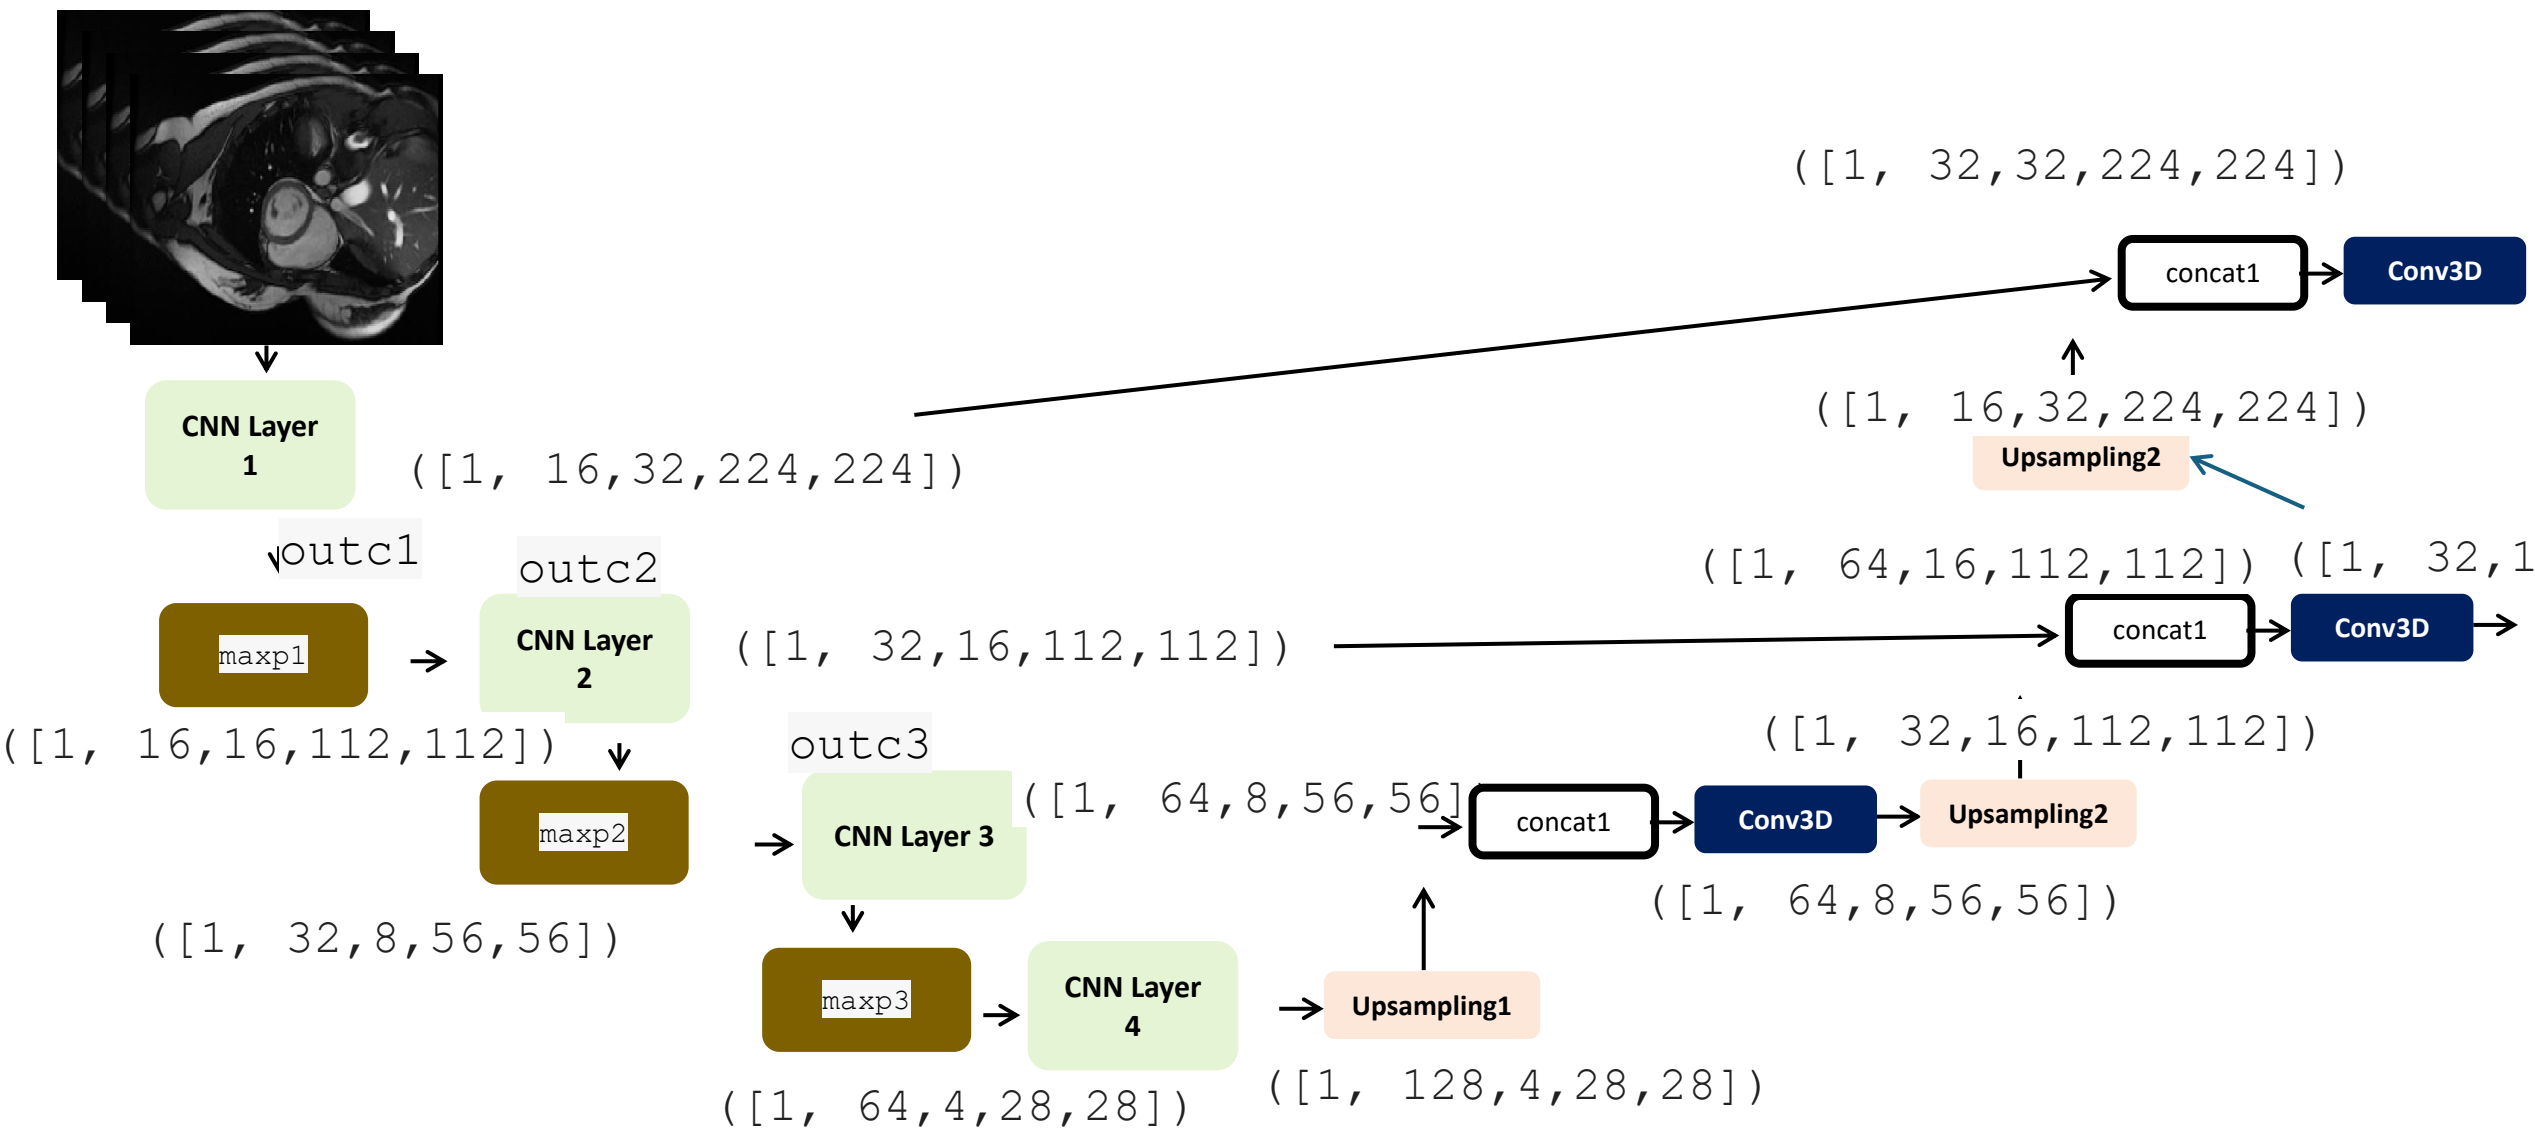
$([1, 32, 16, 112, 112])$

$([1, 64, 16, 112, 112])$

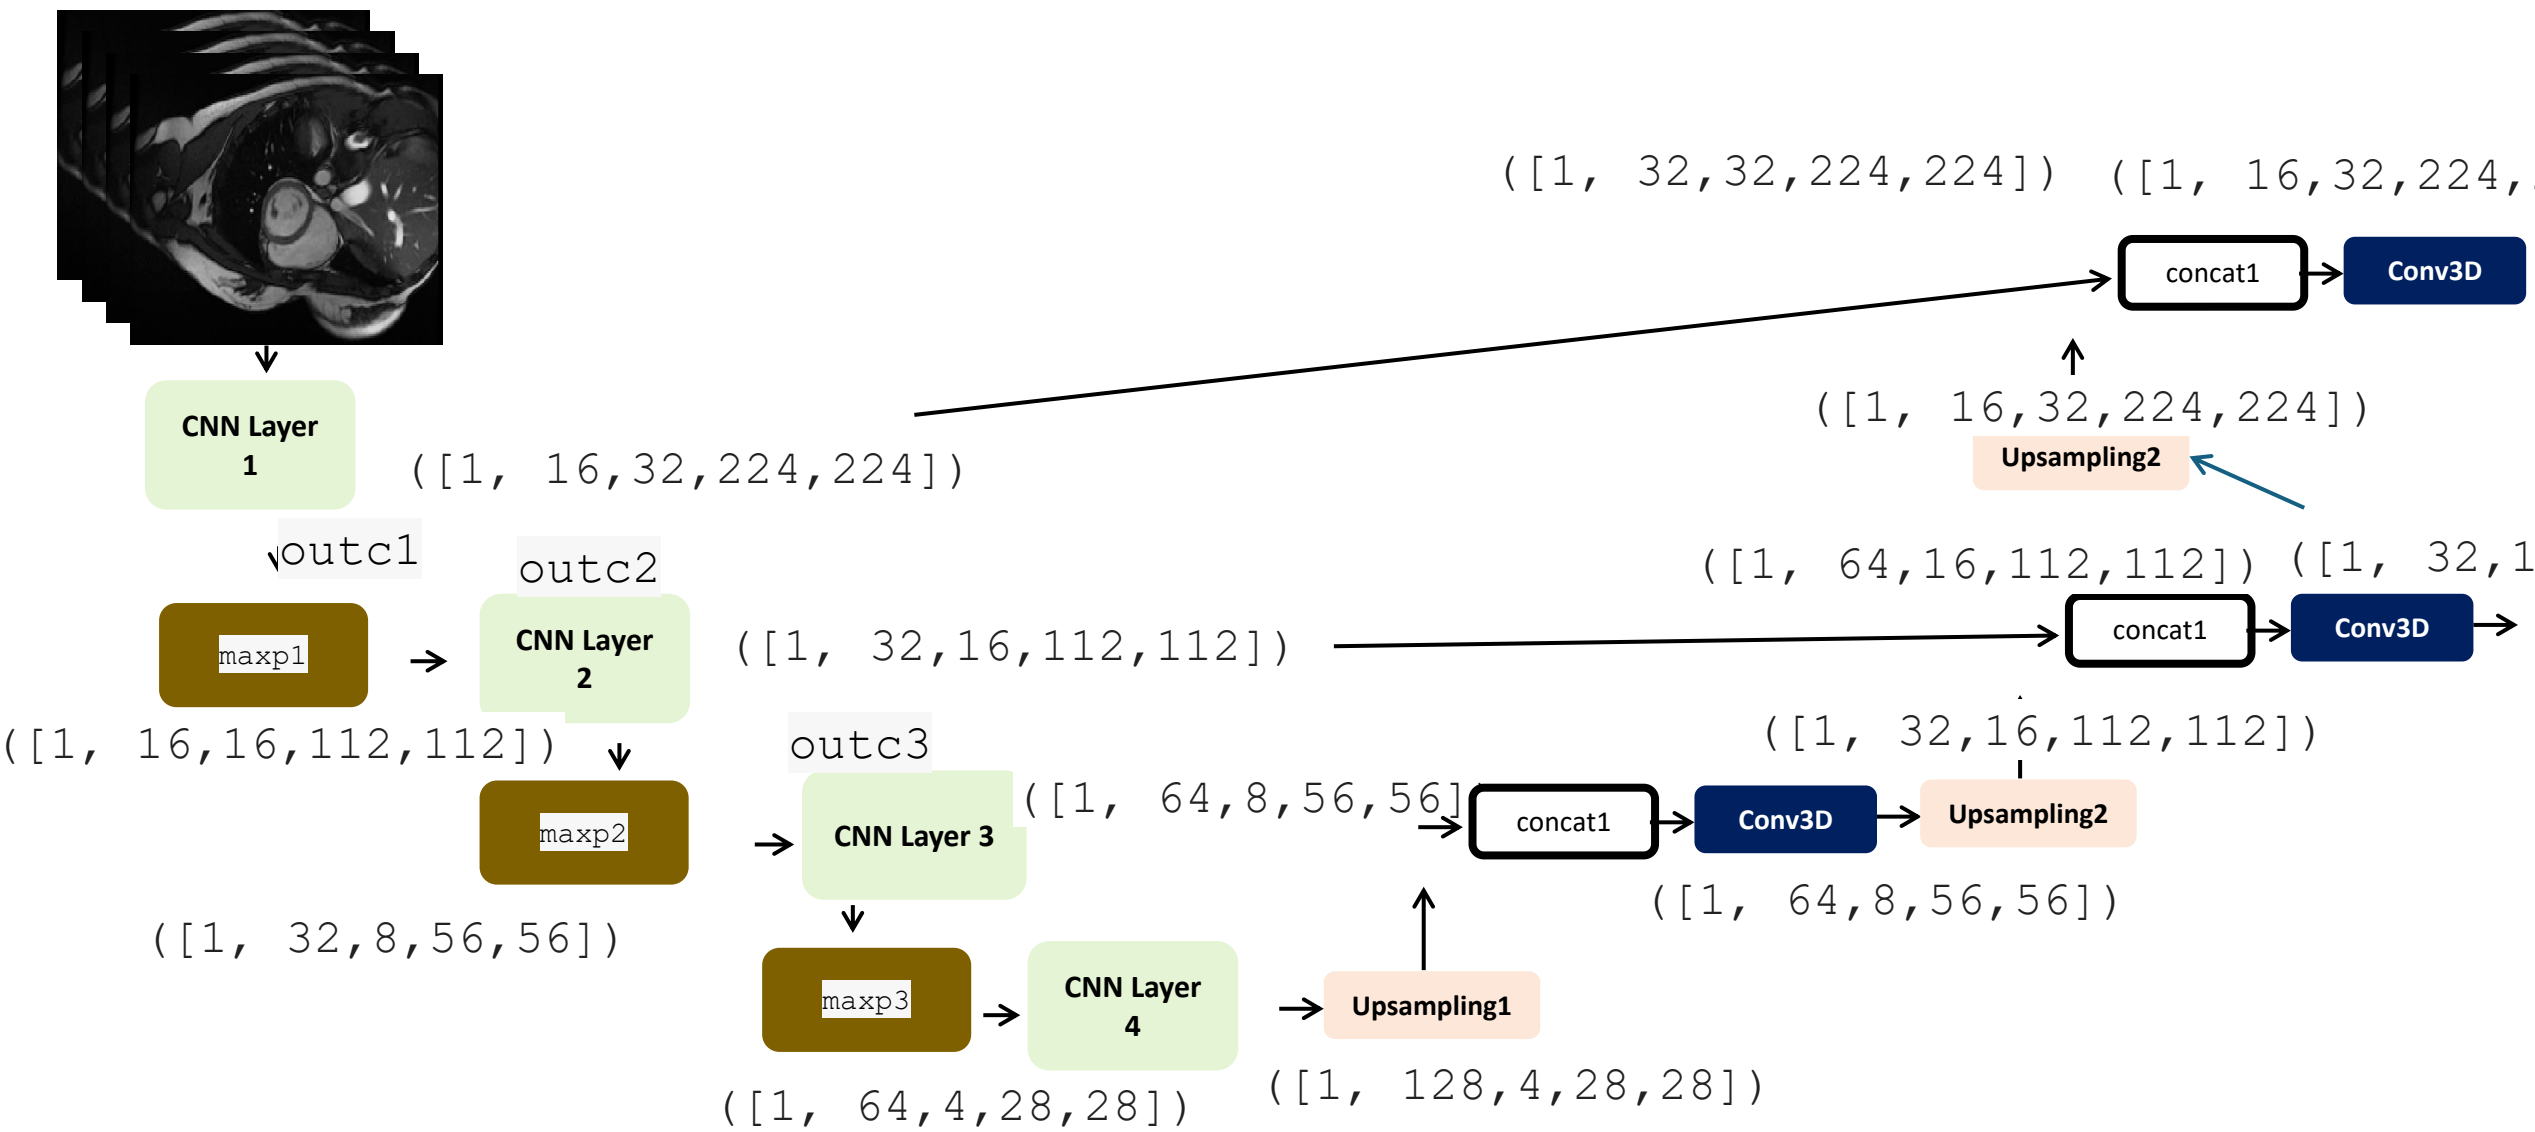
$([1, 32, 16, 112, 112])$



$1 \times 1 \times 132 \times 224 \times 224$



$1 \times 1 \times 132 \times 224 \times 224$



$1 \times 1 \times 132 \times 224 \times 224$

